

Deepfake Audio Detection: Real vs. Synthetically Cloned Voices

AI-600 Deep Learning — Midterm Report

Course: AI-600 Deep Learning, Spring 2026 **Submitted:** 08 May 2026

Name	Student ID
Ahmad Gill	25280065
Imran Saeed	25280083
Israr Ali	25280102

Abstract

The proliferation of neural voice cloning systems has made it trivially easy to synthesize speech perceptually indistinguishable from a target speaker, enabling large-scale voice fraud, political misinformation, and authentication bypass at scale. Recent evaluations confirm that LLM-based text-to-speech architectures now produce synthetic audio that defeats systems effective against conventional vocoders [G], underscoring the urgency of this problem. This project develops a binary classifier to distinguish genuine human speech from synthetically generated audio using the ASVspoof 2019 Logical Access (LA) benchmark [2], a dataset of 121,000+ utterances spanning 19 TTS and voice conversion attack types. Our approach builds on end-to-end architectures [4][5] and self-supervised feature representations from wav2vec 2.0 [7][8], augmented with RawBoost data augmentation [9] to improve generalization to unseen spoofing attacks — the core challenge identified across the benchmark literature [10]. We have completed a thorough exploratory analysis of the ASVspoof 2019 LA dataset, confirming its split structure, class imbalance, and speaker-disjoint properties. Implementation of the RawNet2 baseline is underway, with full training and AASIST integration planned for the remainder of the project.

1. Introduction

The deepfake audio threat has evolved from a theoretical concern to a practical attack vector. Singh et al. [G] evaluate detection systems against LLM-based TTS architectures — including Dia2 and Mayal — and find that models effective against conventional vocoder-based synthesis fail significantly when confronted with newer LLM-based generation methods. Voice-cloning services now operate commercially at low cost, enabling impersonation attacks on voice-authenticated banking systems, injection of synthetic speech into judicial proceedings, and automated generation

of political misinformation at scale. This threat landscape creates an urgent demand for robust, generalizable countermeasures that extend beyond controlled laboratory conditions.

The ASVspoof challenge series [2][3] has driven systematic benchmarking of spoofing countermeasures since 2015, with the 2019 edition introducing 19 distinct TTS and voice conversion attack types across the Logical Access (LA) track. The LA benchmark is deliberately structured to stress-test generalization: the training and development partitions expose only six known attacks (A01–A06), while the evaluation partition introduces 13 previously unseen attacks (A07–A19) [2]. This split directly simulates the real-world scenario where detectors must remain effective against novel synthesis methods encountered post-deployment. Müller et al. [10] formalize the severity of this challenge: models achieving sub-5% Equal Error Rate (EER) on the ASVspoof 2019 LA evaluation set suffer performance degradation of up to 1,000% when evaluated on real-world deepfake recordings of public figures, demonstrating that benchmark performance and real-world generalization remain fundamentally decoupled.

This project directly addresses the generalization challenge by implementing and evaluating a progression of detection architectures on the ASVspoof 2019 LA benchmark. Beginning with the RawNet2 end-to-end baseline [4], the project advances to the AASIST graph-attention architecture [5] augmented with self-supervised feature representations from wav2vec 2.0 [7], following evidence that fine-tuned SSL front-ends substantially improve cross-condition robustness [8]. RawBoost data augmentation [9] is applied to simulate telephony channel effects and real-world noise corruption. System performance is measured using EER and tandem Detection Cost Function (t-DCF) [3] on the LA evaluation partition. Section 2 reviews the relevant literature across four paradigms. Section 3 describes the proposed methodology. Section 4 details the experimental setup. Section 5 presents current progress and preliminary findings. Section 6 outlines the remaining plan, and Section 7 concludes.

2. Literature Review

2.1 Datasets and Evaluation Benchmarks

The ASVspoof challenge has served as the primary driver of anti-spoofing research since its inception in 2015, with successive editions escalating the diversity and realism of synthetic attacks. The ASVspoof 2019 LA dataset [2] contains 121,000+ utterances from 107 speakers drawn from the VCTK corpus, with spoofed speech generated by 19 TTS and voice conversion systems (A01–A19). The dataset employs a speaker-disjoint design: no speaker appears in more than one partition, ensuring that models cannot exploit speaker identity as a shortcut to correct classification. The training set contains 2,580 bonafide and 22,800 spoofed utterances; the development set contains 2,548 bonafide and 22,296 spoofed; and the evaluation set contains 7,355 bonafide and 63,882

spoofed utterances [2].

The ASVspoof 2019 challenge [3] established two official evaluation metrics: the Equal Error Rate (EER), defined as the operating threshold θ^* at which the false alarm rate equals the false rejection rate:

$$\text{FAR}(\theta^*) = \text{FRR}(\theta^*)$$

and the tandem Detection Cost Function (t-DCF), which measures countermeasure cost within a complete speaker verification pipeline by weighting detection errors by their relative impact on end-to-end system performance [3]. Both metrics are lower-is-better, and evaluation is always reported on the evaluation partition rather than development. The official challenge baselines are CQCC+GMM (EER: 9.57%, min t-DCF: 0.2366) and LFCC+GMM (EER: 8.09%, min t-DCF: 0.2116) [3], which 27 of 48 participating teams were able to surpass.

2.2 Feature Extraction Approaches

Early spoofing countermeasures relied on handcrafted spectral features — Mel-Frequency Cepstral Coefficients (MFCC), Linear Frequency Cepstral Coefficients (LFCC), and Constant-Q Cepstral Coefficients (CQCC) — which capture vocoder artefacts through frequency domain analysis. While computationally efficient, these representations generalize poorly: Müller et al. [10] show that CQT-based features outperform mel-spectrograms by approximately 37% relative EER on in-the-wild data, and LFCC-based systems collapse catastrophically on out-of-domain recordings, degrading from 2.98% to 29.42% EER between ASVspoof 2019 and 2015 evaluation conditions [8].

Self-supervised learning (SSL) representations have emerged as a more generalizable alternative. Baevski et al. [7] introduce wav2vec 2.0, a framework that masks speech input in the latent space of a 7-block CNN encoder and learns contextualized representations through contrastive learning over quantized latents across a 12- or 24-block Transformer. Pre-trained on up to 60,000 hours of unlabelled speech, wav2vec 2.0 produces layered representations where shallow layers retain acoustic detail and deeper layers encode phonetic and semantic content [7]. Wang and Yamagishi [8] establish that when a fine-tuned wav2vec 2.0 XLSR front-end is combined with a deep LSTM-based backend, the system achieves 0.11% EER on ASVspoof 2019 LA — crucially, fine-tuning the SSL front-end end-to-end is essential; frozen representations yield 1.47% EER under the same conditions [8]. A comprehensive layer-wise analysis [C] further confirms that intermediate transformer layers carry significantly more anti-spoofing discriminative information than the final layer, with the optimal extraction depth varying by SSL architecture.

2.3 Detection Architectures

The detection architecture literature has evolved from CNN-based spectrogram classifiers toward end-to-end and graph-structured models. Tak et al. [4] introduced the first end-to-end anti-spoofing system that processes raw waveforms directly through a bank of fixed sinc-parameterized filters — deliberately kept non-learnable to avoid overfitting to the six known training attacks — followed by six residual blocks with Feature Map Scaling (FMS) attention and a GRU layer for temporal modelling. The RawNet2 single system achieves 4.66% EER on the ASVspoof 2019 LA evaluation set; score-level fusion with an LFCC-GMM baseline improves this to 1.12% EER [4].

AASIST [5] advances the architecture to a heterogeneous graph neural network. A RawNet2-style encoder feeds spectral and temporal artefact representations into separate graph branches processed by stacked Heterogeneous-Stacking Graph Attention Layers (HS-GALs). A max graph operation introduces competitive selection among graph nodes, and a stack node aggregates cross-domain artefacts before a fully-connected scoring head. Even the lightweight AASIST-L variant, with only 85K parameters, achieves 0.99% EER, outperforming all competing single systems at the time of publication [5]. Wang and Yamagishi [6] provide a systematic multi-seed comparison of neural countermeasures, finding that intra-model EER variance across random seeds can span 1.2–3.1% on the same configuration — confirming that single-run evaluation substantially underestimates result reliability. The 2026 resolution-aware detection framework [F] achieves 0.16% EER through multi-resolution spectral alignment with cross-scale attention, establishing the current performance frontier on this benchmark.

2.4 Generalization and Robustness

The dominant finding of the deepfake audio detection literature is a profound generalization gap between controlled benchmark performance and real-world detection capability. Müller et al. [10] evaluate 12 architectures on ASVspoof 2019 and a newly collected 37.9-hour In-the-Wild (ITW) dataset of celebrity deepfakes sourced from social media. The best-performing ASVspoof system (RawGAT-ST, 1.23% EER) degrades to 37.15% EER on ITW data — a degradation of over 1,000%. The authors demonstrate that adding ITW data to training does not substantially improve ITW performance (33.94% $\rightarrow$ 33.14% EER), indicating that the problem is distributional mismatch rather than data quantity.

Recent 2025 work has begun to systematically address this gap. A generalizable artifact-focused approach [A] proposes explicit Artifact Detection Modules (ADMs) that disentangle synthesis artifacts from speaker identity, achieving an F1 score of 0.813 on the ITW dataset. Lopez et al. [B] demonstrate that WavLM and Wav2Vec2 backbones combined with novel augmentation and loss functions surpass the top single system in the ASVspoof 5 Challenge. Nguyen-Le et al. [I] introduce AFSS (Artifact-Focused Self-Synthesis), which generates pseudo-fake samples at training time to enforce artifact-centric learning, achieving a mean EER of 5.45% across seven evaluation datasets

— the strongest published generalization result as of early 2026. These findings motivate our use of RawBoost augmentation [9] and SSL feature representations [7][8] as core design choices targeting robustness rather than in-distribution benchmark performance.

Building on these findings, our proposed methodology integrates state-of-the-art end-to-end detection with SSL-based feature extraction and targeted data augmentation, with a particular focus on cross-attack generalization on the ASVspoof 2019 LA evaluation set.

3. Methodology

3.1 Problem Formulation

Deepfake audio detection is formulated as a binary classification problem. Given an input waveform $\mathbf{x}$ of variable length, a detection function $f : \mathbf{x} \rightarrow s$ maps the input to a real-valued detection score s , where lower scores indicate higher probability of genuine (bonafide) speech and higher scores indicate synthetic (spoofed) speech. A decision threshold θ determines the binary classification output:

$$\hat{y} = \begin{cases} 0 & \text{(bonafide) if } s < \theta \\ 1 & \text{(spoof) if } s \geq \theta \end{cases}$$

The threshold θ^* is chosen to minimize EER during evaluation. Unlike speaker verification, no enrollment utterance is required — the classifier operates solely on the test waveform’s acoustic properties.

3.2 Feature Extraction

The methodology investigates two complementary feature extraction pipelines operating in parallel:

End-to-End Stream. Raw waveforms at 16 kHz are processed directly by a bank of sinc-parameterized filters, following the RawNet2 design [4]. These filters are initialized with Mel-scale frequency spacing and kept fixed (non-learnable) to prevent overfitting to the frequency characteristics of the six known training attacks — a design choice explicitly motivated by generalization concerns [4].

SSL Stream. Raw waveforms are additionally processed through a pre-trained wav2vec 2.0 encoder [7]. Following the empirical findings of Wang and Yamagishi [8] and the layer-wise analysis of [C], features are extracted from intermediate transformer layers rather than the final layer, as intermediate representations retain more acoustically grounded anti-spoofing information. The SSL encoder is fine-tuned end-to-end with the detection backend; frozen representations are used only as

an ablation baseline, consistent with the finding that fine-tuning is essential for strong performance [8].

All input waveforms are processed as 4-second (64,000 sample at 16 kHz) segments. Variable-length utterances are padded with silence or randomly cropped to this fixed length during training. The pipeline can be summarized as follows:

Raw Waveform (16 kHz)

→ Sinc Filters (fixed) → 6 Residual Blocks (FMS) → GRU(1024)

→ wav2vec 2.0 (intermediate layers, fine-tuned)

Heterogeneous Graph (AASIST HS-GAL)

Max Graph Pool → Readout ←

Detection Score (s)

3.3 Model Architecture

Baseline: RawNet2 [4]. RawNet2 processes raw waveforms through a bank of sinc filters parameterized by fixed lower and upper cutoff frequencies arranged on the Mel scale. Six residual blocks equipped with Feature Map Scaling (FMS) attention follow, where FMS applies filter-wise scaling to emphasize informative frequency bands. A GRU layer with 1024 units captures temporal dependencies before a fully-connected output layer produces a two-class detection score. The full system processes 64,000-sample (4-second) waveform segments and achieves 4.66% EER on the ASVspoof 2019 LA evaluation set as a single system [4].

Proposed: AASIST [5]. AASIST extends the RawNet2 encoder with a heterogeneous graph neural network designed to jointly model spectral and temporal artefacts. Encoder features feed into two graph branches — a spectral graph G_s and a temporal graph G_t — each processed by stacked HS-GAL layers. A max graph operation (MGO) introduces competitive node selection across the two branches, and a stack node acts as a classification token aggregating cross-domain evidence. The model is trained with random graph pooling (50% spectral, 30% temporal node removal) as implicit regularization [5]. Even the lightweight AASIST-L variant with 85K parameters achieves 0.99% EER on the LA evaluation set, making it a practical deployment target.

3.4 Data Augmentation

To improve robustness to channel effects and real-world acoustic degradation, RawBoost [9] augmentation is applied during training. RawBoost operates directly on raw waveforms without requiring external noise recordings or impulse response databases, using three stochastic noise components:

1. **Linear and non-linear convolutive noise (LnL-CM):** Models transmission channel effects using Hammerstein systems with five notch filters spanning 20–8,000 Hz.
2. **Impulsive signal-dependent additive noise (ISD-ADD):** Simulates burst noise with a relative power ratio in the range 0–10%.
3. **Stationary signal-independent additive noise (SSI-ADD):** Adds FIR-coloured Gaussian noise at 10–40 dB SNR.

Tak et al. [9] demonstrate that combining modes + achieves a 44% reduction in EER and a 27% relative reduction in min t-DCF on ASVspoof 2021 LA, outperforming SpecAugment, WavAugment, and codec augmentation. The noise-aware robustness analysis of [E] further motivates augmentation-based training: state-of-the-art detectors suffer significant EER degradation under noisy conditions (0–10 dB SNR), with fine-tuning reducing this gap by 10–15 percentage points.

Training Configuration. Models are trained using weighted binary cross-entropy to address class imbalance (approximately 9:1 spoof-to-bonafide ratio in the training set). The Adam optimizer is used with an initial learning rate of 10^{-4} . For AASIST, cosine annealing is applied following the original training protocol [5]. Mini-batches of 32 utterances are used, with each waveform randomly cropped to 64,000 samples. Training runs for up to 100 epochs with early stopping based on development EER (patience: 5 epochs). All experiments are repeated with three random seeds, and results are reported as mean $\pm$ standard deviation following the multi-seed protocol recommended by [6].

4. Experimental Setup

4.1 Dataset

All experiments use the ASVspoof 2019 Logical Access (LA) track [2][3]. The dataset is downloaded from the Edinburgh DataShare repository (doi:10.7488/ds/2461) and organized into three speaker-disjoint partitions. Our exploratory data analysis confirmed the following statistics from the protocol files:

Split	Bonafide	Spoofed	Total	Speakers	Known Attacks
Train	2,580	22,800	25,380	20	A01–A06
Development	2,548	22,296	24,844	20	A01–A06

Split	Bonafide	Spoofed	Total	Speakers	Known Attacks
Evaluation	7,355	63,882	71,237	67	A07–A19

Speaker overlap analysis confirmed zero intersection between all three partitions, ensuring the speaker-disjoint property holds. The class imbalance is consistent across splits, with spoofed utterances comprising approximately 89.8% of each partition — a characteristic that motivates the use of weighted loss functions.

The evaluation partition includes attacks A07–A19, none of which appear in training or development. This design makes performance on the evaluation set a direct proxy for cross-attack generalization. Attack A17 (VAE-based voice conversion) and A10 (Tacotron2 + WaveRNN) have been identified as the hardest evaluation attacks in both the database paper [2] and challenge overview [3].

4.2 Evaluation Metrics

The primary evaluation metric is the Equal Error Rate (EER), computed at the detection threshold where the false alarm rate and false rejection rate are equal. The secondary metric is the minimum tandem Detection Cost Function (min t-DCF), which accounts for the cost of countermeasure errors within a full ASV pipeline [3]. Both metrics are lower-is-better, and all results are reported on the LA evaluation set rather than the development set.

4.3 Implementation Details

All models are implemented in PyTorch with torchaudio for audio loading. RawBoost augmentation is applied on-the-fly during training as a waveform-level transform. Official code repositories for RawNet2 and AASIST serve as reference implementations for architecture verification. Experiments run on available GPU hardware. The ASVspooof 2019 LA protocol files and .flac audio files are verified to be complete and consistent with the published dataset statistics prior to model training.

5. Preliminary Results and Current Progress

5.1 Completed Work

The following milestones have been achieved as of the midterm submission date:

- **Dataset acquisition and verification.** The full ASVspooof 2019 LA dataset has been downloaded, organized, and verified. Protocol files for all three splits have been parsed and confirmed to be complete. Exploratory data analysis has been completed covering split

size distribution, label distribution, speaker overlap, attack type distribution, and audio file availability.

- **Exploratory data analysis.** A detailed EDA notebook documents all dataset properties relevant to model design. The class imbalance (89.8% spoof across all splits) has been confirmed and will be addressed through weighted cross-entropy. The speaker-disjoint property across all three partitions has been verified programmatically (zero speaker overlap in all pairwise comparisons). Audio file counts were verified against protocol row counts, confirming dataset integrity.
- **Architecture review and implementation planning.** All ten core reference papers have been read, summarized, and integrated into a literature review. Architecture components for RawNet2 and AASIST have been documented with reference to official implementations. Training hyperparameters have been fixed following the original papers.

5.2 Baseline Results

Model training is currently in progress. The table below reports published baseline numbers against which our reproductions will be compared. Our results will be added as training completes.

Model	Year	EER (%)	min t-DCF	Notes
CQCC + GMM [3]	2019	9.57	0.2366	Official challenge baseline
LFCC + GMM [3]	2019	8.09	0.2116	Official challenge baseline
RawNet2 [4]	2021	4.66	0.1294	Single system (linear sinc)
AASIST [5]	2022	0.83	0.0275	Single system
AASIST-L [5]	2022	0.99	0.0309	Lightweight variant (85K params)
Resolution- Aware [F]	2026	0.16	—	Current SOTA
Our RawNet2 (reprod.)	2026	—	—	In progress

Model	Year	EER (%)	min t-DCF	Notes
Our AASIST + SSL (WIP)	2026	—	—	Planned

5.3 EDA Findings Relevant to Modeling

The exploratory analysis revealed several dataset properties that directly inform modeling decisions:

1. **Class imbalance is severe and consistent.** With 89.8% spoof samples across all splits, unweighted cross-entropy would produce a biased classifier. Weighted CE or OC-Softmax will be used.
2. **Speaker disjoint design is confirmed.** Zero speaker overlap across train, development, and evaluation partitions ensures that speaker identity cannot be exploited as a detection shortcut.
3. **Known-unknown attack split.** Training attacks A01–A06 include two voice conversion and four TTS systems. The evaluation introduces 13 entirely new attack types — the generalization stress test is built into the benchmark protocol.
4. **Dataset is complete and verified.** All audio files referenced in the protocol files are present and accounted for.

6. Future Work and Remaining Plan

6.1 Immediate Next Steps (Next Two Weeks)

The immediate priority is completing the RawNet2 baseline reproduction. Training will proceed using the hyperparameters documented in Section 3.4, with results validated against the published 4.66% EER single-system result [4]. Upon confirming reproducibility within an acceptable margin ($\pm 0.5\%$ EER), AASIST training will begin using the official implementation as a reference.

Concurrently, the SSL feature extraction pipeline will be implemented. wav2vec 2.0 features will be extracted from intermediate transformer layers following the layer-wise analysis in [C], with the extraction depth determined by a validation sweep over the development set. Fine-tuning will be performed end-to-end following the protocol established by [8], which yielded a 10–20x improvement over frozen representations.

RawBoost augmentation will be integrated and ablated: training runs with no augmentation, mode only, mode only, and combined + will be compared to quantify the augmentation contribution, matching the experimental design of [9].

6.2 Planned Experiments

The following experiments constitute the core ablation study for the final report:

- **Feature $\times$ Architecture grid.** All combinations of (raw sinc, SSL extracted, SSL fine-tuned) $\times$ (RawNet2, AASIST) will be evaluated, yielding six baseline configurations.
- **Augmentation ablation.** Each RawBoost mode will be evaluated independently and in combination with the strongest architecture.
- **Per-attack EER breakdown.** Individual EER will be reported for each of the 13 evaluation attacks (A07–A19) to identify which attack types pose the greatest challenge and whether SSL features improve performance on the hardest attacks (A10, A17).
- **Cross-dataset evaluation.** Following the generalization evaluation protocol of [10], the strongest trained model will be evaluated on the In-the-Wild dataset to measure the generalization gap.
- **Noise robustness evaluation.** Motivated by [E], detection performance will be measured under controlled additive noise at SNR levels of 0, 5, 10, and 20 dB to characterize robustness beyond clean benchmark conditions.

6.3 Final Report Targets

The final report will present a complete ablation table across all feature/architecture/augmentation combinations, accompanied by per-attack analysis and cross-dataset evaluation results. The project repository will be submitted with fully reproducible training code, pre-trained model weights, and evaluation scripts. Comparison against the resolution-aware 2026 SOTA [F] will contextualize the final system’s performance relative to the current frontier. Emerging paradigms — including artifact-focused self-synthesis [I] and large audio-language model reasoning [H] — will be discussed as directions for future investigation beyond this project scope.

7. Conclusion

This work addresses the critical challenge of detecting synthetically generated speech in the face of increasingly sophisticated voice cloning technologies, including LLM-based TTS systems that are now evaluated against existing detectors [G]. Using the ASVspoof 2019 Logical Access benchmark [2], the project implements a structured comparison of end-to-end and SSL-augmented detection architectures, with RawBoost augmentation [9] targeting the generalization gap documented by [10]. As of the midterm, the dataset has been fully acquired and verified through comprehensive exploratory analysis, architecture implementations are underway, and the experimental plan is fully specified. The completed system will produce a systematic ablation study identifying which combinations of features, architectures, and augmentation strategies are most critical to cross-attack

generalization — the central open challenge in deepfake audio detection research.

References

- [1] M. Li, Y. Ahmadiadli, and X.-P. Zhang, “A Survey on Speech Deepfake Detection,” *arXiv:2404.13914*, 2024 (updated Jul 2025).
- [2] X. Wang et al., “ASVspoof 2019: A large-scale public database of synthesized, converted and replayed speech,” *Computer Speech and Language*, vol. 64, 2020.
- [3] M. Todisco et al., “ASVspoof 2019: Future Horizons in Spoofed and Fake Audio Detection,” in *Proc. Interspeech*, 2019.
- [4] H. Tak, J. Patino, M. Todisco, A. Nautsch, N. Evans, and A. Larcher, “End-to-End Anti-Spoofing with RawNet2,” in *Proc. ICASSP*, 2021, pp. 6369–6373.
- [5] J.-W. Jung et al., “AASIST: Audio Anti-Spoofing using Integrated Spectro-Temporal Graph Attention Networks,” in *Proc. ICASSP*, 2022.
- [6] X. Wang and J. Yamagishi, “A Comparative Study on Recent Neural Spoofing Countermeasures for Synthetic Speech Detection,” in *Proc. Interspeech*, 2021.
- [7] A. Baevski, H. Zhou, A. Mohamed, and M. Auli, “wav2vec 2.0: A Framework for Self-Supervised Learning of Speech Representations,” in *Proc. NeurIPS*, 2020.
- [8] X. Wang and J. Yamagishi, “Investigating Self-Supervised Front Ends for Speech Spoofing Countermeasures,” in *Proc. Odyssey*, 2022.
- [9] H. Tak, M. Kamble, J. Patino, M. Todisco, and N. Evans, “RawBoost: A Raw Data Boosting and Augmentation Method Applied to Automatic Speaker Verification Anti-Spoofing,” in *Proc. ICASSP*, 2022, pp. 6382–6386.
- [10] N. Müller et al., “Does Audio Deepfake Detection Generalize?” in *Proc. Interspeech*, 2022.
- [11] “Beyond Identity: A Generalizable Approach for Deepfake Audio Detection,” *arXiv:2505.06766*, 2025.
- [12] J. A. Lopez, G. Stemmer, and H. C. Maruri, “Generalizable Detection of Audio Deepfakes,” *arXiv:2507.01750*, 2025.
- [13] “Comprehensive Layer-wise Analysis of SSL Models for Audio Deepfake Detection,” *arXiv:2502.03559*, 2025.
- [14] U. Sen, A. Luqman, and A. Chattopadhyay, “Toward Noise-Aware Audio Deepfake Detection,” *arXiv:2512.13744*, 2025.

- [15] K. A. Shahriar et al., “Lightweight Resolution-Aware Audio Deepfake Detection via Cross-Scale Attention and Consistency Learning,” *arXiv:2601.06560*, 2026.
- [16] R. Singh et al., “Audio Deepfake Detection in the Age of Advanced Text-to-Speech Models,” *arXiv:2601.20510*, 2026.
- [17] Korzh et al., “Towards Robust Speech Deepfake Detection via Human-Inspired Reasoning,” *arXiv:2603.10725*, 2026.
- [18] Nguyen-Le et al., “AFSS: Artifact-Focused Self-Synthesis for Mitigating Bias in Audio Deepfake Detection,” *arXiv:2603.26856*, 2026.